

N-Grams and Corpus Linguistics

Julia Hirschberg

CS 4705

Linguistics vs. Engineering

- “But it must be recognized that the notion of “probability of a sentence” is an entirely useless one, under any known interpretation of this term.”

Noam Chomsky (1969)

- “Anytime a linguist leaves the group the recognition rate goes up.”

Fred Jelinek† (1988)



THE 2008 CAMPAIGN: The Message and Corporate Marketing

The Words They Used

The words that the speakers have used during the Democratic convention suggest how the party's themes have changed since the last presidential campaign.

Speakers have hammered home Barack Obama's "change" theme, using the word about eight times as often as they

did in 2004.

Also, unlike 2004, when the Kerry campaign sought to avoid direct attacks on the president at the convention, the speakers have regularly have been mentioning John McCain by name. Speakers in 2004 practiced "the art of the


implicit slam," a veteran Democratic speechwriter said then, indirectly bashing Mr. Bush while barely using his name.

Also on the upswing: more mentions of the economy, energy, Iran and Iraq.

Words less frequently used: freedom, Sept. 11 and terrorism.

2008


DENVER



Source: Federal News Service transcripts

2004

BOSTON



MATTHEW ERICSON/THE NEW YORK TIMES

Next Word Prediction

- From a NY Times story...
 - Stocks ...
 - Stocks plunged this
 - Stocks plunged this morning, despite a cut in interest rates
 - Stocks plunged this morning, despite a cut in interest rates by the Federal Reserve, as Wall ...
 - Stocks plunged this morning, despite a cut in interest rates by the Federal Reserve, as Wall Street began

- Stocks plunged this morning, despite a cut in interest rates by the Federal Reserve, as Wall Street began trading for the first time since last ...
- Stocks plunged this morning, despite a cut in interest rates by the Federal Reserve, as Wall Street began trading for the first time since last Tuesday's terrorist attacks.

Human Word Prediction

- Clearly, at least some of us have the ability to predict future words in an utterance.
- How?
 - Domain knowledge: red blood vs. red hat
 - Syntactic knowledge: the...<adj|noun>
 - Lexical knowledge: baked <potato vs. steak>

Claim

- A useful part of the knowledge needed to allow **Word Prediction** can be captured using simple statistical techniques
- In particular, we'll be interested in the notion of the **probability** of a sequence (of letters, words,...)

Useful Applications

- Why do we want to predict a word, given some preceding words?
 - Rank the **likelihood** of sequences containing various alternative hypotheses, e.g. for ASR

Theatre owners say popcorn/unicorn sales have doubled...

- Assess the likelihood/goodness of a sentence, e.g. for text generation or machine translation

The doctor recommended a cat scan.

El doctor recomendó una exploración del gato.

N-Gram Models of Language

- Use the previous $N-1$ words in a sequence to predict the next word
- Language Model (LM)
 - unigrams, bigrams, trigrams,...
- How do we **train** these models?
 - Very large corpora

Corpora

- Corpora are online collections of text and speech
 - Brown Corpus
 - Wall Street Journal
 - AP newswire
 - Hansards
 - Timit
 - DARPA/NIST text/speech corpora (Call Home, Call Friend, ATIS, Switchboard, Broadcast News, Broadcast Conversation, TDT, Communicator)
 - TRAINS, Boston Radio News Corpus

Counting Words in Corpora

- What is a word?
 - e.g., are **cat** and **cats** the same word?
 - **September** and **Sept**?
 - **zero** and **oh**?
 - Is **_** a word? ***** ?) . ,
 - How many words are there in **don' t** ? **Gonna** ?
 - In Japanese and Chinese text -- how do we identify a word?

Terminology

- **Sentence**: unit of written language
- **Utterance**: unit of spoken language
- **Word Form**: the inflected form as it actually appears in the corpus
- **Lemma**: an abstract form, shared by word forms having the same **stem**, part of speech, word sense – stands for the class of words with same **stem**
- **Types**: number of distinct words in a corpus (vocabulary size)
- **Tokens**: total number of words

Simple N-Grams

- Assume a language has T word types in its lexicon, how likely is word x to follow word y ?
 - Simplest model of word probability: $1/T$
 - **Alternative 1**: estimate likelihood of x occurring in new text based on its general frequency of occurrence estimated from a corpus (**unigram** probability)
popcorn is more likely to occur than **unicorn**
 - **Alternative 2**: condition the likelihood of x occurring in the context of previous words (**bigrams**, **trigrams**,...)
mythical unicorn is more likely than **mythical popcorn**

Computing the Probability of a Word Sequence

- Compute the product of component conditional probabilities?

- $P(\text{the mythical unicorn}) = P(\text{the}) * P(\text{mythical}|\text{the}) * P(\text{unicorn}|\text{the mythical})$

- But...the *longer* the sequence, the *less likely* we are to find it in a training corpus

$P(\text{Most biologists and folklore specialists believe that in fact the mythical unicorn horns derived from the narwhal})$

- What can we do?

Bigram Model

- Approximate $P(w_n|w_1^{n-1})$ by $P(w_n|w_{n-1})$
 - E.g., $P(\text{unicorn}|\text{the mythical})$ by $P(\text{unicorn}|\text{mythical})$
- **Markov assumption**: the probability of a word depends only on the probability of a limited history
- **Generalization: the probability of a word depends only on the probability of the n previous words**
 - trigrams, 4-grams, 5-grams,...
 - the higher n is, the more data needed to train
 - **backoff** models...

Using N-Grams

- For N-gram models

- $P(w_n | w_1^{n-1}) \approx P(w_n | w_{n-N+1}^{n-1})$

- E.g. for bigrams, $P(w_8 | w_1^{8-1}) \approx P(w_8 | w_{8-2+1}^{8-1})$

- $P(w_{n-1}, w_n) = P(w_n | w_{n-1}) P(w_{n-1})$

- $P(w_{8-1}, w_8) = P(w_8 | w_7) P(w_7)$

- By the Chain Rule we can decompose a joint probability, e.g. $P(w_1, w_2, w_3)$ as follows

$$P(w_1, w_2, \dots, w_n) = P(w_1 | w_2, w_3, \dots, w_n) P(w_2 | w_3, \dots, w_n) \dots P(w_{n-1} | w_n) P(w_n)$$

$$P(w_1^n) \approx \prod_{k=1}^n P(w_k | w_1^{k-1})$$

- For bigrams then, the probability of a sequence is just the product of the conditional probabilities of its bigrams, e.g.

$$P(\text{the, mythical, unicorn}) = P(\text{unicorn}|\text{mythical}) \\ P(\text{mythical}|\text{the}) P(\text{the}|\text{<start>})$$

Training and Testing

- N-Gram probabilities come from a **training corpus**
 - overly narrow corpus: probabilities don't **generalize**
 - overly general corpus: probabilities don't **reflect task or domain**
- A separate **test corpus** is used to **evaluate** the model
 - held out test set; development (dev) test set
 - Simple baseline
 - results tested for statistical significance – how do they differ from a baseline? Other results?

A Simple Bigram Example

- Estimate the likelihood of the sentence **I want to eat Chinese food.**
 - $P(\text{I want to eat Chinese food}) = P(\text{I} \mid \langle \text{start} \rangle) P(\text{want} \mid \text{I}) P(\text{to} \mid \text{want}) P(\text{eat} \mid \text{to}) P(\text{Chinese} \mid \text{eat}) P(\text{food} \mid \text{Chinese}) P(\langle \text{end} \rangle \mid \text{food})$
- What do we need to calculate these likelihoods?
 - Bigram probabilities for each word pair sequence in the sentence
 - Calculated from a large corpus

Early Bigram Probabilities from BERP

Eat on	.16	Eat Thai	.03
Eat some	.06	Eat breakfast	.03
Eat lunch	.06	Eat in	.02
Eat dinner	.05	Eat Chinese	.02
Eat at	.04	Eat Mexican	.02
Eat a	.04	Eat tomorrow	.01
Eat Indian	.04	Eat dessert	.007
Eat today	.03	Eat British	.001

<start> I	.25	Want some	.04
<start> I' d	.06	Want Thai	.01
<start> Tell	.04	To eat	.26
<start> I' m	.02	To have	.14
I want	.32	To spend	.09
I would	.29	To be	.02
I don' t	.08	British food	.60
I have	.04	British restaurant	.15
Want to	.65	British cuisine	.01
Want a	.05	British lunch	.01

- $P(\text{I want to eat British food}) = P(\text{I}|\text{<start>}) P(\text{want}|\text{I}) P(\text{to}|\text{want}) P(\text{eat}|\text{to}) P(\text{British}|\text{eat}) P(\text{food}|\text{British}) = .25 * .32 * .65 * .26 * .001 * .60 = .000080$
 - Suppose $P(\text{<end>}|\text{food}) = .2$?
 - How would we calculate **I want to eat Chinese food** ?
- Probabilities roughly capture ``syntactic" facts and ``world knowledge"
 - **eat** is often followed by an NP
 - British food is not too popular
- N-gram models can be trained by **counting** and **normalization**

Early BERP Bigram Counts

	I	Want	To	Eat	Chinese	Food	lunch
I	8	1087	0	13	0	0	0
Want	3	0	786	0	6	8	6
To	3	0	10	860	3	0	12
Eat	0	0	2	0	19	2	52
Chinese	2	0	0	0	0	120	1
Food	19	0	17	0	0	0	0
Lunch	4	0	0	0	0	1	0

Early BERP Bigram Probabilities

- Normalization: divide each row's counts by appropriate unigram counts for w_{n-1}

I	Want	To	Eat	Chinese	Food	Lunch
3437	1215	3256	938	213	1506	459

- Computing the bigram probability of **I I**
 - $C(I,I)/C(I \text{ in call contexts})$
 - $p(I|I) = 8 / 3437 = .0023$
- **Maximum Likelihood Estimation** (MLE): relative frequency
$$\frac{freq(w_1, w_2)}{freq(w_1)}$$

What do we learn about the language?

- What's being captured with ...
 - $P(\text{want} \mid I) = .32$
 - $P(\text{to} \mid \text{want}) = .65$
 - $P(\text{eat} \mid \text{to}) = .26$
 - $P(\text{food} \mid \text{Chinese}) = .56$
 - $P(\text{lunch} \mid \text{eat}) = .055$
- What about...
 - $P(I \mid I) = .0023$
 - $P(I \mid \text{want}) = .0025$
 - $P(I \mid \text{food}) = .013$

- $P(I \mid I) = .0023$ I I I I want
- $P(I \mid \text{want}) = .0025$ I want I want
- $P(I \mid \text{food}) = .013$ the kind of food I want is ...

Ngrams vs. FSAs

- How is Ngram modeling stronger than FSAs?
- For what kinds of tasks would each be better?
- What about automatic speech recognition (ASR)?

Approximating Shakespeare

- Generating sentences with random unigrams...
 - Every enter now severally so, let
 - Hill he late speaks; or! a more to leg less first you enter
- With bigrams...
 - What means, sir. I confess she? then all sorts, he is trim, captain.
 - Why dost stand forth thy canopy, forsooth; he is this palpable hit the King Henry.
- Trigrams
 - Sweet prince, Falstaff shall die.
 - This shall forbid it should be branded, if renown made it empty.

- Quadrigrams
 - What! I will go seek the traitor Gloucester.
 - Will you not tell me who I am?
 - What's coming out here looks like Shakespeare because it *is* Shakespeare
- Note: *As we increase the value of N , the accuracy of an n -gram model increases, since choice of next word becomes increasingly constrained*

N-Gram Training Sensitivity

- If we repeated the Shakespeare experiment but trained our n-grams on a Wall Street Journal corpus, what would we get?
- Note: *This question has major implications for corpus selection or design*

WSJ is *not* Shakespeare: Sentences Generated from WSJ

unigram: Months the my and issue of year foreign new exchange's september were recession exchange new endorsed a acquire to six executives

bigram: Last December through the way to preserve the Hudson corporation N. B. E. C. Taylor would seem to complete the major central planners one point five percent of U. S. E. has already old M. X. corporation of living on information such as more frequently fishing to keep her

trigram: They also point to ninety nine point six billion dollars from two hundred four oh six three percent of the rates of interest stores as Mexico and Brazil on market conditions

Evaluation and Data Sparsity Questions

- **Perplexity** and **entropy**: how do you *estimate* how well your language model fits a corpus once you're done?
- **Smoothing and Backoff**: how do you handle unseen n-grams?

Perplexity and Entropy

- Information theoretic metrics
 - Useful in measuring how well a **grammar** or **language model (LM)** models a natural language or a corpus
- **Entropy**: With 2 LMs and a corpus, which LM is the better match for the corpus? How much information is there (in e.g. a grammar or LM) about what the next word will be? More is better!
 - For a random variable X ranging over e.g. bigrams and a probability function $p(x)$, the entropy of X is the expected negative log probability
$$H(X) = - \sum_{x=1}^{x=n} p(x) \log_2 p(x)$$

- Entropy is the lower bound on the # of bits it takes to encode information e.g. about bigram likelihood

- Cross Entropy

- An upper bound on entropy derived from estimating true entropy by a subset of possible strings – we don't know the real probability distribution

- Perplexity $PP(W) = 2^{H(W)}$

- At each choice point in a grammar
 - What are the average number of choices that can be made, weighted by their probabilities of occurrence?
 - I.e., Weighted average branching factor
- How much probability does a grammar or language model (LM) assign to the sentences of a corpus, compared to another LM? The more information, the lower perplexity

Some Useful Observations

- There are 884,647 tokens, with 29,066 word form types, in an approximately one million word Shakespeare corpus
 - Shakespeare produced 300,000 bigram types out of 844 million possible bigrams: so, *99.96% of the possible bigrams were never seen (have zero entries in the table)*
- A small number of events occur with high frequency
- A large number of events occur with low frequency
- You can quickly collect statistics on the high frequency events
- You might have to wait an arbitrarily long time to get valid statistics on low frequency events
- Some zeroes in the table are really zeros But others are simply low frequency events you haven't seen yet. How to address?


Smoothing

- Words follow a Zipfian distribution
 - Small number of words occur very frequently
 - A large number are seen only once
 - Zipf's law: *a word's frequency is approximately inversely proportional to its rank in the word distribution list*
- Zero probabilities on one bigram cause a zero probability on the entire sentence
- So....how do we estimate the likelihood of unseen n-grams?

Smoothing is like Robin Hood: Steal from the rich and give to the poor (in probability mass)


- We often want to make predictions from sparse statistics:

$P(w \mid \text{denied the})$
3 allegations
2 reports
1 claims
1 request
7 total



- Smoothing flattens spiky distributions so they generalize better

$P(w \mid \text{denied the})$
2.5 allegations
1.5 reports
0.5 claims
0.5 request
2 other
7 total



- Very important all over NLP, but easy to do badly!

Laplace Smoothing

- For unigrams:
 - Add 1 to every word (type) count to get an adjusted count c^*
 - Normalize by N (#tokens) + V (#types)
 - Original unigram probability
 - New unigram probability

$$P(w_i) = \frac{c_i}{N}$$

$$P_w(w_i) = \frac{c_i + 1}{N + V}$$

Unigram Smoothing Example

- Tiny Corpus, V=4; N=20

$$P_{LP}(w_i) = \frac{c_i + 1}{N + V}$$

Word	True Ct	Unigram Prob	New Ct	Adjusted Prob
eat	10	.5	11	.46
British	4	.2	5	.21
food	6	.3	7	.29
happily	0	.0	1	.04
	20	1.0	~20	1.0

- *So, we lower some (larger) observed counts in order to include unobserved vocabulary*

- For bigrams:

- Original
$$P(w_n | w_{n-1}) = \frac{c(w_n | w_{n-1})}{c(w_{n-1})}$$

- New
$$P(w_n | w_{n-1}) = \frac{c(w_n | w_{n-1}) + 1}{c(w_{n-1}) + V}$$

- *But this change counts drastically:*

- *Too much weight* given to unseen ngrams
- In practice, unsmoothed bigrams often work better!
- Can we smooth more usefully?

Good-Turing Discounting

- Re-estimate amount of probability mass for zero (or low count) ngrams by looking at ngrams with higher counts
 - Estimate
$$c^* = (c + 1) \frac{N_{c+1}}{N_c}$$
 - E.g. N_0 's adjusted count is a function of the count of ngrams that occur once, N_1
 - Assumes:
 - Word bigrams each follow a binomial distribution
 - We know number of unseen bigrams ($V \times V$ -seen)

Backoff Methods (e.g. Katz '87)

- For e.g. a trigram model
 - Compute unigram, bigram and trigram probabilities
 - Usage:
 - Where trigram unavailable **back off** to bigram if available, o.w. **back off** to the current word's unigram probability
 - E.g **An omnivorous unicorn**
- ***NB: check errata pages for changes to figures***

Class-Based Backoff

- **Back off** to the class rather than the word
 - Particularly useful for proper nouns (e.g., names)
 - Use count for the number of names in place of the particular name
 - E.g. **< N | friendly >** instead of **< dog | friendly >**

Smoothing Summed Up

- Add-one smoothing (easy, but inaccurate)
 - Add 1 to every word count (Note: this is type)
 - Increment normalization factor by Vocabulary size: N (tokens) + V (*types*)
- Good-Turing
 - Re-estimate amount of probability mass for zero (or low count) ngrams by looking at ngrams with higher counts
- Backoff models
 - When a count for an n-gram is 0, back off to the count for the (n-1)-gram
 - These can be weighted – trigrams count more
- Class-based smoothing
 - For certain types of n-grams, back off to the count of its syntactic class....semantic category??
 - E.g., Count ProperNouns in place of names (e.g., Obama)

Google NGrams

All Our N-gram are Belong to You

By Peter Norvig - 8/03/2006 11:26:00 AM

Posted by Alex Franz and Thorsten Brants, Google Machine Translation Team

Here at Google Research we have been using word [n-gram models](#) for a variety of R&D projects, such as [statistical machine translation](#), speech recognition, [spelling correction](#), entity detection, information extraction, and others. While such models have usually been estimated from training to share this enormous dataset with everyone. We processed 1,024,908,267,229 words of running text and are publishing the counts for all 1,176,470,663 five-word sequences that appear at least 40 times. There are 13,588,391 unique words, after discarding words that appear less than 200 times.

Example

- serve as the incoming 92
- serve as the incubator 99
- serve as the independent 794
- serve as the index 223
- serve as the indication 72
- serve as the indicator 120
- serve as the indicators 45
- serve as the indispensable 111
- serve as the indispensable 40
- serve as the individual 234

Summary

- N-gram probabilities can be used to *estimate* the likelihood
 - Of a word occurring in a context (N-1)
 - Of a sentence occurring at all
- Entropy and perplexity can be used to evaluate the information content of a language and the goodness of fit of a LM or grammar
- Smoothing techniques and backoff models deal with problems of unseen words in corpus
- Next Class: Read Ch. 5 on parts-of-speech (POS)